



Who Makes What?

Prediction Models of Who Makes Over 50k

Wintana Tadesse

Jon Weisbecker

Alireza Esfandiari



Purpose & motivation

- ❖ To make a predictive model of whether or not someone makes over 50k
- Market segmentation
- Characteristics of the two groups



Data collection.

- Source of dataset: the US Census
- Data contains 32561 observations
- Includes one target variable and 14 explanatory variables

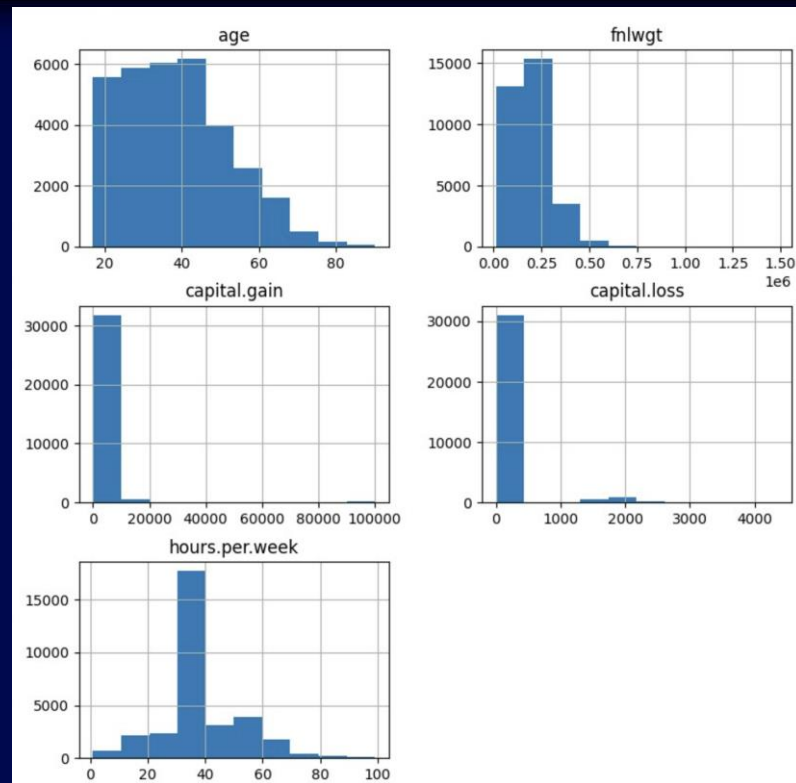



Dataset and null data

#	Column	Non-Null Count		Dtype
0	age	32561	non-null	int64
1	workclass	32561	non-null	object
2	fnlwgt	32561	non-null	int64
3	education	32561	non-null	object
4	education.num	32561	non-null	int64
5	marital.status	32561	non-null	object
6	occupation	32561	non-null	object
7	relationship	32561	non-null	object
8	race	32561	non-null	object
9	sex	32561	non-null	object
10	capital.gain	32561	non-null	int64
11	capital.loss	32561	non-null	int64
12	hours.per.week	32561	non-null	int64
13	native.country	32561	non-null	object
14	income	32561	non-null	object




Dataset...





Data Split




we splitted our data into 20% training and 80% test set with random state = 200 because our data set is large.


Previous Models

****As we are going to predict a class variable we chose these models****


Decision Tree





KNN





SVM



New Models



- ❖ Logistic Regression
- ❖ DNN
- ❖ Autoencoder
- ❖ GNB





Models/Results

Parameters

```
from sklearn.neighbors import KNeighborsClassifier  
model_KNN = KNeighborsClassifier()  
parameters_KNN = {'n_neighbors' : range(1,20)}
```

Results

```
Fitting 5 folds for each of 19 candidates, totalling 95 fits  
KNeighborsClassifier {'n_neighbors': 19} 0.7999845166843597 0.7941040994933211 0.7452889439434973
```

Parameters

```
from sklearn.naive_bayes import GaussianNB
model_GNB = GaussianNB()
parameters_GNB = {'var_smoothing': np.logspace(0,-9, num=100)}
```

Results

Fitting 5 folds for each of 100 candidates, totalling 500 fits

GaussianNB {'var_smoothing': 1.2328467394420658e-05} 0.805129038070459 0.7976354982343006 0.7598007329112451



Logistic Regression

Parameters

```
from sklearn.linear_model import LogisticRegression
model_LR = LogisticRegression()
parameters_LR = {'C' : np.logspace(-2, 2, 10)}
```

Results

```
LogisticRegression {'C': 0.0774263682681127} 0.7987945343379923 0.7913403961308153 0.7507386774102396
```




Decision Tree

Parameters

```
from sklearn.tree import DecisionTreeClassifier
model_DT = DecisionTreeClassifier()
parameters_DT = {'criterion' : ['gini', 'entropy'],
                  'max_leaf_nodes' : [2, 10, 100],
                  'min_samples_split' : [2, 10, 100],
                  'max_depth': list(range(3,10)),
                  'min_samples_leaf' : [1, 2]
                  }
```

Results

Fitting 5 folds for each of 252 candidates, totalling 1260 fits

DecisionTreeClassifier {'criterion': 'gini', 'max_depth': 9, 'max_leaf_nodes': 100, 'min_samples_leaf': 1, 'min_samples_split': 2} 0.8596824262156633 0.8550591125441425 0.8477944986756594

Parameters

```
## DNN model
DNN_model = Sequential()
DNN_model.add(Dense(units = 16, input_dim = X_train.shape[1], activation='relu'))
DNN_model.add(Dense(units = 8, activation='relu'))
DNN_model.add(Dense(units = 8, activation='relu'))
DNN_model.add(Dense(1, activation='sigmoid'))

print(DNN_model.summary())
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 16)	1600
dense_1 (Dense)	(None, 8)	136
dense_2 (Dense)	(None, 8)	72
dense_3 (Dense)	(None, 1)	9

=====
Total params: 1,817
Trainable params: 1,817
Non-trainable params: 0

```
epochs = 10, batch_size = 32, verbose=1,
```

Results


loss: 0.5625 - accuracy: 0.7620 - f1_score: 0.0348

Train

val_loss: 0.5547 - val_accuracy: 0.7634 - val_f1_score: 0.0280

Test

DNN



Autoencoder


Parameters

```
# Adding a classifier layer
clf_input_layer = Input(shape=(encoding_dim,))
clf_layer = Dense(1, activation='sigmoid')(clf_input_layer) #This can be any layer such as RNN, CNN, ...
classifier = Model(inputs=clf_input_layer, outputs=clf_layer)

# Compile the classifier
classifier.compile(optimizer=Adam(learning_rate=0.001), loss='binary_crossentropy', metrics=['accuracy'])

# Train the classifier
model_fit = classifier.fit(encoded_X_train, y_train, epochs=50, batch_size=32, shuffle=True,
                           validation_data=(encoded_X_test, y_test), verbose=0)
```

Results



Classifier Test Loss: 6.6469
Classifier Test Accuracy: 0.3204



Feature Selection (Variance)

```
Variance values:
  age                1.860614e+02
  fnlwgt             1.114080e+10
  capital.gain       5.454254e+07
  capital.loss       1.623769e+05
  hours.per.week     1.524590e+02
  ...
  native.country_Thailand      5.525199e-04
  native.country_Trinidad&Tobago 5.831976e-04
  native.country_United-States  9.330010e-02
  native.country_Vietnam        2.053505e-03
  native.country_Yugoslavia     4.911590e-04
  Length: 99, dtype: float64
```



Feature Selection Logistic Regression

Parameters

```
# Apply variance threshold feature selection  
selector = VarianceThreshold(threshold=0.1)
```

```
clf = LogisticRegression(random_state = 0)
```

Results

	precision	recall	f1-score	support
False	0.81	0.94	0.87	4945
True	0.62	0.29	0.39	1568
accuracy			0.79	6513
macro avg	0.71	0.62	0.63	6513
weighted avg	0.76	0.79	0.76	6513

```
Index(['age', 'fnlwgt', 'capital.gain', 'capital.loss', 'hours.per.week',  
      'workclass_Private', 'education_Bachelors', 'education_HS-grad',  
      'education_Some-college', 'marital.status_Married-civ-spouse',  
      'marital.status_Never-married', 'occupation_Adm-clerical',  
      'occupation_Craft-repair', 'occupation_Exec-managerial',  
      'occupation_Prof-specialty', 'relationship_Not-in-family',  
      'relationship_Own-child', 'race_White', 'sex_Male'],
```



Feature Selection KNN

Parameters

```
# Apply variance threshold feature selection  
selector = VarianceThreshold(threshold=0.1)
```

```
clf = KNeighborsClassifier(n_neighbors = 19)
```

Results

	precision	recall	f1-score	support
False	0.80	0.98	0.88	4945
True	0.77	0.21	0.33	1568
accuracy			0.79	6513
macro avg	0.78	0.59	0.60	6513
weighted avg	0.79	0.79	0.75	6513

```
Index(['age', 'fnlwgt', 'capital.gain', 'capital.loss', 'hours.per.week',  
      'workclass_Private', 'education_Bachelors', 'education_HS-grad',  
      'education_Some-college', 'marital.status_Married-civ-spouse',  
      'marital.status_Never-married', 'occupation_Adm-clerical',  
      'occupation_Craft-repair', 'occupation_Exec-managerial',  
      'occupation_Prof-specialty', 'relationship_Not-in-family',  
      'relationship_Own-child', 'race_White', 'sex_Male'],  
      dtype=object)
```



Feature Selection GNB

Parameters

```
# Apply variance threshold feature selection  
selector = VarianceThreshold(threshold=0.1)
```

```
clf = GaussianNB(var_smoothing = 1.2328467394420658e-05)
```

Results

	precision	recall	f1-score	support
False	0.81	0.97	0.88	4945
True	0.72	0.26	0.38	1568
accuracy			0.80	6513
macro avg	0.76	0.61	0.63	6513
weighted avg	0.78	0.80	0.76	6513

```
Index(['age', 'fnlwgt', 'capital.gain', 'capital.loss', 'hours.per.week',  
      'workclass_Private', 'education_Bachelors', 'education_HS-grad',  
      'education_Some-college', 'marital.status_Married-civ-spouse',  
      'marital.status_Never-married', 'occupation_Adm-clerical',  
      'occupation_Craft-repair', 'occupation_Exec-managerial',  
      'occupation_Prof-specialty', 'relationship_Not-in-family',  
      'relationship_Own-child', 'race_White', 'sex_Male'],
```




Feature Selection Decision Tree

Parameters

```
# Apply variance threshold feature selection  
selector = VarianceThreshold(threshold=0.1)
```

```
clf = DecisionTreeClassifier(max_depth = 9, max_leaf_nodes = 100, min_samples_leaf = 1, min_samples_split = 2)
```

Results

	precision	recall	f1-score	support
False	0.88	0.94	0.91	4945
True	0.76	0.58	0.66	1568
accuracy			0.85	6513
macro avg	0.82	0.76	0.78	6513
weighted avg	0.85	0.85	0.85	6513

```
Index(['age', 'fnlwgt', 'capital.gain', 'capital.loss', 'hours.per.week',  
      'workclass_Private', 'education_Bachelors', 'education_HS-grad',  
      'education_Some-college', 'marital.status_Married-civ-spouse',  
      'marital.status_Never-married', 'occupation_Adm-clerical',  
      'occupation_Craft-repair', 'occupation_Exec-managerial',  
      'occupation_Prof-specialty', 'relationship_Not-in-family',  
      'relationship_Own-child', 'race_White', 'sex_Male'],
```



Feature Selection Vs. Usual (KNN- DT- GNB- LR)

Model	Accuracy	F1-Score
DT	0.859	0.847
GNB	0.80	0.76
KNN	0.794	0.745
LogisticRegression	0.798	0.75

Usual

Model	Accuracy	F1-Score
DT	0.85	0.85
GNB	0.80	0.76
KNN	0.79	0.75
LogisticRegression	0.79	0.76

Feature Selection



Final Comparison

Performance Measures

Model	Accuracy	F1-Score	Best Model
DT	0.859	0.847	★
DNN	0.763	0.028	
Autoencoder	0.32		
GNB	0.797	0.759	
KNN	0.794	0.745	
SVM	0.79	0.88	
LogisticRegression	0.798	0.75	



Conclusion

- Based on both graph and statistical comparison, Decision tree model is still more fitted model Vs all the other models.

PS: however, it may be better to choose the DT with feature selection as it uses less feature hence less resources.

- As shown in the comparison, An autoencoder model was the lowest performer. Autoencoder model is a neural network model that can be used to learn a compressed representation of raw data and input values, however our data set is large and have lots of details so maybe if we reduce the number of input values, we can make the model less likely be confused by tiny (and irrelevant) details.



Recommendation

- Usually we use neural networks when we do forecasting and time series applications, sentiment analysis and other text applications. It is not recommend for studies like this one where we have a binary output because:
 - Hard to interpret most of the times
 - They require too much data
 - They take time to be developed
 - They take a lot of time in the training phase



THANK YOU!!

Q&A